

Amazon Review Opinion Search Engine – Final Report

Alexander Yue

COSC 4397

Arjun Mukherjee

Summer 2025

1. Overview

This project implements an opinion-aware retrieval system for customer reviews using a real-world Amazon dataset. The system supports structured aspect-opinion queries and outputs review IDs most relevant to the query semantics and sentiment. Three models were developed:

- **Baseline:** Boolean matching
- **Method 1:** Boolean + rating-aligned sentiment filtering
- **Method 2:** Sentence-BERT + semantic similarity + rating alignment

Evaluation is based on precision of retrieval across five fixed queries:

- *Audio Quality:Poor*
- *Wifi Signal:Strong*
- *Mouse Button:Click Problem*
- *GPS Map:Useful*
- *Image Quality:Sharp*

2. Background

In classical IR systems, exact keyword matches via Boolean operators limit the system's ability to understand nuanced human language. This is especially problematic for opinion retrieval, where context, polarity, and paraphrasing all influence relevance.

Given a dataset of **210,761** reviews, this project explores layered strategies:

- *Lexical matching - Matching based on exact words or phrases in the text.*
- *Heuristic sentiment scoring - Assigning sentiment (positive/negative/neutral) to text using rule-based or statistical techniques.*
- *Semantic embeddings - Mapping words, sentences, or documents into high-dimensional vector space where meaning is captured.*

The reviews are preprocessed and stored as a Pickle file (reviews_segment.pkl) for efficient retrieval.

3. Query Format

All queries follow this format:

[aspect: opinion]

Where:

- aspect: max 2 tokens (e.g., "audio quality")
- opinion: 1–2 tokens (e.g., "poor", "click problem")

Fixed benchmark queries used:

- audio quality: poor
- wifi signal: strong
- mouse button: click problem
- gps map: useful
- image quality: sharp

4. Methodology

4.1 Baseline – Boolean Matching

Implementation:

- Reviews are cleaned using stopword removal and lowercasing.
- A review is marked as a match if the cleaned review text contains both the aspect and opinion strings.

Strengths:

- Fast, rule-based
- Simple to implement

Limitations:

- Fails on semantically similar rewordings
 - ex: For “Click Problem”, some reviews mention clicking, for instance one review mentions “right mouse button becomes stuck and doesn’t click at all”
 - This review does not get caught by Boolean Matching
- No sentiment disambiguation
- Prone to false positives from subword containment

4.2 Method 1 – Boolean Matching + Rating Filter

Enhancement: Add sentiment polarity filtering using review star ratings.

Pipeline Logic:

1. Execute Boolean match as in baseline.
2. Determine if the query opinion is positive or negative.
3. Filter matched reviews:
 - Positive opinion → keep if rating > 3
 - Negative opinion → keep if rating ≤ 3

Advantages:

- Aligns opinion sentiment with actual review rating
- Reduces sentiment-inverted matches

Limitations:

- Relies on predefined opinion lexicon
- Still limited by lexical overlap

4.3 Method 2 – Sentence-BERT + Semantic Filtering + Rating

Architecture:

- Uses sentence-transformers with the all-MiniLM-L6-v2 model.
- Embeds both query and review texts using Sentence-BERT.
- Computes cosine similarity; filters with a threshold of 0.6.
- Applies rating filter just like in Method 1.

Technical Details:

- Embedding dimension: 384
 - Each sentence or text input is converted into a 384-dimensional vector.
- Cosine similarity via `util.cos_sim``
 - a function that measures how **similar** two vectors are by the **angle** between them.
- Batch inference supported with GPU acceleration
 - Grouping multiple inputs together to process them simultaneously, improving efficiency and speed. Used for speed optimization with CUDA GPU's

Advantages:

- Captures paraphrasing and semantic similarity

- Outperforms prior methods in nuanced queries

Challenges:

- High compute/memory cost
 - ~830s for this method compared to ~85s for method1
- Threshold tuning required
 - Thresholds tested:
 - .6, .3, .45, .7

5. Evaluation Strategy

Each method was evaluated on its ability to retrieve relevant review IDs for the five predefined aspect-opinion queries. The primary metric used for comparison was precision, defined as:

$$\text{Precision} = (\text{Relevant Matches}) / (\text{Total Matches Returned})$$

Relevant reviews were identified by copying queried reviews into an LLM, and asking which ones were relevant to each request. Example request:

- “Out of the following reviews, which one is relevant to Audio Quality:Poor”

Because the accuracy of LLM’s are not 100%, there may be missed or false positive reviews, particularly for larger review sets. In fact, during review of the results, some false positive/negative reviews were identified and corrected. While there is a likely a high number of false positive/negatives that were missed, since relevant reviews will be graded by an LLM, and it would be near impossible to assess all results by hand, it was decided this was a good choice for finding relevant reviews

6. Results Summary

Query	Baseline(Boolean) Time Taken: 82.40S			Method1 (Boolean + Rating Filter) Time Taken: 86.94 seconds			Method2 (Semantic Similarity with BERT) Time Taken: 830.57		
	#Ret.	#Rel	Prec.	#Ret.	#Rel	Prec.	#Ret.	#Rel	Prec.
Audio quality:poor	40	14	.35	23	20	.8696	7	6	.8571
Wifi signal:strong	5	3	.6	4	4	1	2	2	1
Mouse button:click problem	1	0	0	0	0	n/a	9	7	.7778
Gps map:useful	5	4	.80	5	4	.80	154	92	.5794
Image quality:sharp	103	5	.0485	82	2	.0244	11	7	.6363

Tests were run on a machine with the specs:

- OS: Windows 10 Pro
- AMD Ryzen 5 1600, 6 Cores, 3700Mhz clock
- 24GB Vram
- NVIDIA RTX 2080 6GB Vram

Discussion:

- **Baseline** often returned many irrelevant results due to shallow matching.
- **Method 1** showed improved filtering for sentiment alignment, especially for queries like “poor audio quality.”
- **Method 2** returned reviews that were exactly relevant, rather than keyword searching that boolean applies.

- **Semantic Models Matter:** Queries like "gps map: useful" retrieved diverse phrasing like "helpful for navigation" or "perfect trip companion" under Method 2 but were missed by Boolean methods.
- **False Positives in Baseline:** For "wifi signal: strong", the baseline incorrectly matched "weak signal" due to only checking for "signal".
- **Rating Filtering Reduces Noise:** Several irrelevant matches were removed in Method 1 and Method 2 through sentiment-aware filtering.

8. Conclusion

This project demonstrates how progressively sophisticated retrieval techniques improve results and relevance of opinion queries. While the baseline Boolean technique is viable as a simple, quick, easy to implement starting point, semantic models such as Sentence-BERT meaningfully boosts precision, as well as identifying reviews potentially missed by baseline boolean models. However, it comes at a massive computational, and time cost, as Baseline Boolean techniques all took ~80-90s on CPU, while Semantic Models took ~850+ seconds on a CUDA core system. Ultimately, this project highlights the power of transformer-based sentence encoders, proving their usefulness in real-world IR tasks, as well as showing the benefits and value of hybrid pipelines that combine semantic similarity, with sentiment alignment heuristics.