

Project Walkthrough

COSC 4397

Presented by Ziming Zhao and Yang Lu

Summer 2025

Overview

- Amazon review opinion search engine implementation
- Project description:
https://www2.cs.uh.edu/~arjun/courses/nlp_ugrad/hw_proj/res_proj/res_proj.pdf
- 40% of your total credit!
- Submittable:
 - a PDF report,
 - a folder named “codes” with one “readme.txt” for detailed instructions on how to run the code,
 - a folder named “outputs” with your generated outputs.
 - Compress all files into a zip file. Name is lastname_PID.zip, e.g. smith_1234.zip
- Due date: 8/1/2025

In brief

- Input format (Query): 2-word aspect + opinion
 - Aspect is always two words
 - Opinion is always 1 or 2 words
 - Connotation is required, e.g “The sound was incredibly muffled and lacked any clarity”
- Output – related review ids in a .txt file
- Dataset
 - Source: http://www2.cs.uh.edu/%7Earjun/courses/nlp_ugrad/hw_proj/res_proj/res_proj.7z
 - Total reviews: 210,761
 - Rating distribution: 1 ★ (28611), 2 ★ (15561), 3 ★ (19224), 4 ★ (38709), 5 ★ (108656)

Step 1 – Implement the baseline: Boolean Search

- Will substring containment check with Python in statement work?

No, think about subwords: **weak** in **tweak**. Think about regex or other SQL based implementation.

- Query results are reviews that contain at least one of the aspect words.
Query examples that may work:
 - “audio quality: poor”, the return reviews should include at least “audio” or “quality” or both of them
 - (audio AND poor) OR (quality AND poor)
 - (audio OR quality) AND poor
 - “mouse button: click problem”, (mouse AND click AND problem) OR (button AND click AND problem)

Step 2 – Implement two advanced methods

- Compare with the baseline using precision as evaluation metrics
- Turn in the evaluation result of Baseline & Method1 & Method 2 in the report with the following table:

Query	Baseline (Boolean)			Method 1 (M1)			Method 2 (M2)		
	# Ret.	# Rel.	Prec.	# Ret.	# Rel.	Prec.	# Ret.	# Rel.	Prec.
audio quality:poor									
wifi signal:strong									
mouse button:click problem									
gps map:useful									
image quality:sharp									

Tip

Baseline & M1 wont be graded. Only M2 (which should be the best performed method) will be graded.

Project Evaluation: 50% Report, 50% Performance Judge

- We will only evaluate the retrieval result from your best method M2 (so try implement and test various SOTAs for better result).
- We will match your result with the ground truth based on 4 tests as below (from grading scheme)

Tests 1 to 3 will be matched with the ground truth. For Test 4, we will leverage LLMs to generate ground truth. Each query will be scored over four distinct tests, weighted as follows:

Test 1: Boolean Aspect Term Retrieval (30%) The review must contain at least one aspect word (e.g., audio, quality). No opinion term required.

Test 2: Aspect AND Opinion Match (20%) The review must mention both aspect and opinion terms anywhere in the text. Proximity or sentiment not enforced.

Test 3: Aspect OR Opinion Match (20%) The review may mention either the aspect or the opinion term. Designed to test broader retrieval.

Test 4: Proper Connotation (30%) The review must express the correct sentiment orientation of the opinion term toward the aspect. Preferably in the same sentence.

There will be 5 queries in total. Each query contributes 20 points. Within each query:

- Test 1 (Boolean): 6 points
- Test 2 (AND): 4 points
- Test 3 (OR): 4 points
- Test 4 (Connotation): 6 points

Total: 5 queries \times 20 points = **100 points**

- The `llm_ids` are considered as ground truth for test 4. You will get partial credit based on your

For test 1-3, we will only grade by precision:

$$\text{Precision} = \frac{|\text{input_ids} \cap \text{groundTruth_ids}|}{|\text{input_ids}|}$$

Notice 4.1

You will get partial credits even if your retrieved reviews doesn't exactly match the ground truth, based on your precision.























For test 4, we will compute the f1 score:

$$F1 = \frac{2 \cdot |\text{input_ids} \cap \text{llm_ids}|}{|\text{input_ids}| + |\text{llm_ids}|}$$

Such evaluation will be computed four times on each test based on the points mentioned in Sec 3:

$$\text{test_score} = F1 | \text{precision} \times \# \text{ points}$$

Output Files

-  Baseline model (folder)
 -  audio_quality_test1.txt
 -  audio_quality_test2.txt
 -  audio_quality_test3.txt
 -  wifi_signal_test1.txt
 -  wifi_signal_test2.txt
 -  wifi_signal_test3.txt
 -  mouse_button_test1.txt
 -  mouse_button_test2.txt
 -  mouse_button_test3.txt
 -  gps_map_test1.txt
 -  gps_map_test2.txt
 -  gps_map_test3.txt
 -  image_quality_test1.txt
 -  image_quality_test2.txt
 -  image_quality_test3.txt
-  Advanced models (folder)
 -  audio_quality_test4.txt
 -  wifi_signal_test4.txt
 -  mouse_button_test4.txt
 -  gps_map_test4.txt
 -  image_quality_test4.txt

Step 3 – Format the output

- The output should be in a txt file, with each line representing a review id only.
- The file name should be {aspect}.txt, e.g., audio_quality.txt.
- Use your best model to retrieve the most related reviews!

audio_quality_test1.txt:

```
1 R10019MUX6F9A
2 R1002I943QCT20
3 R1003RILN06MX1
4 R100523NBIQIEV
5 R1006KJEGKGV00
6 R1006WPZ81TXED
7 R1006XNHNIQMZ0
8 R10079U2I4PP1Z
9 R1007LULU4W7YH
10 R10094W7TS9IXU
11 R1009GW4F1WC1B
12 R1009NU0YPYXS7
13 R1009X50E67SI0
14 R100A2D3D7XDJ4
15 R100AERLNTU2HQ
16 R100CNB1MEHAG3
17 R100D2CV4WK16J
```



audio.txt:

```
1 R10019MUX6F9A
2 R1002I943QCT20
3 R1003RILN06MX1
4 R100523NBIQIEV
5 R1006KJEGKGV00
6 R1006WPZ81TXED
7 R1006XNHNIQMZ0
8 R10079U2I4PP1Z
9 R1007LULU4W7YH
10 R10094W7TS9IXU
11 R1009GW4F1WC1B
12 R1009NU0YPYXS7
13 R1009X50E67SI0
14 R100A2D3D7XDJ4
15 R100AERLNTU2HQ
16 R100CNB1MEHAG3
17 R100D2CV4WK16J
```



audio_quality.txt:

```
1 'R10019MUX6F9A'
2 'R1002I943QCT20'
3 'R1003RILN06MX1'
4 'R100523NBIQIEV'
5 'R1006KJEGKGV00'
6 'R1006WPZ81TXED'
7 'R1006XNHNIQMZ0'
8 'R10079U2I4PP1Z'
9 'R1007LULU4W7YH'
10 'R10094W7TS9IXU'
11 'R1009GW4F1WC1B'
12 'R1009NU0YPYXS7'
13 'R1009X50E67SI0'
14 'R100A2D3D7XDJ4'
15 'R100AERLNTU2HQ'
16 'R100CNB1MEHAG3'
17 'R100D2CV4WK16J'
18 'R100DFUV1MA11K'
```



What to submit?

LastName_PID.zip

├ *Codes/*

| ├── *Code1.py*

| ├── *Code2.py*

| └ *README*

├ *Report.pdf*

└ *Outputs/*

├── *audio_quality.txt*

├── *wifi_signal.txt*

├── *mouse_button.txt*

├── *gps_map.txt*

└ *image_quality.txt*

- Name your folder “Codes/” for your code
 - Implementation files
 - README: instructions on how to run your code.
 - If your code for M2 does not run or generate the expected output, you will be penalized.
- Name your folder “Outputs” for your output
 - Only need to submit the retrieval result by your best method M2.
 - You do need to implement both the Baseline and M1 models—it's not possible to evaluate M2 effectively without comparing against them.
- 1 PDF file of your report
 - Name it Report.pdf
- No need to submit any dataset/compiled codes

Project Report Template

- Overview
- Background
- Method
- Design
- Result
- Discussion
- Conclusion

Code Demo: Handle Connotation

We will now go through a simple method named sentence-BERT that deals with connotation.